

Predicting People's Perceived Street Safety with Convolutional Neural Networks

Lingling Zhang
School of Continuing Studies
York University
Toronto, Canada

Durai Nachiappan
School of Continuing Studies
York University
Toronto, Canada

Abstract—This project aimed to build explainable machine learning models that could predict people's perceived safety towards Toronto streets views. We applied Convolutional Neural Network (CNN) techniques on street view images fetched through Google Street View Static API. Our models can classify people's perceived safety towards street views into 2 classes: More Safe and Less Safe. Multiple ensemble learning techniques were attempted and model interpretation tool - LIME was applied to assist us in understanding our models' predictions.

Keywords—Perceived Safety, Street View, Computer Vision, Deep Learning, Convolutional Neural Networks, Transfer Learning, Ensemble Learning, Model Interpretation

I. INTRODUCTION

Understanding people's perceived safety of street view could help us to infer the city scene, provide residents with additional information about where they live, and assist city planners to build a more friendly and equal living environment for the residents. It would also be useful for researchers to examine the social and economic consequences of urban perception.

Our research was built on 2 previous research: the Place Pulse project [1] and the Streetsocore project [2]. The Place Pulse project measured people's perception of safety, class, and uniqueness in four cities - Boston and New York in the United States, and Linz and Salzburg in Austria. In the crowdsourced study, 7,872 participants took part in an online game in which they were shown images using 208,738 pairwise comparisons. The users were shown two images, selected randomly from the dataset and asked to click on one in response to one of the 3 questions: Which place looks safer? Which place looks more upper-class? or Which place looks more unique? The study used 4,136 geo-tagged images, among them the 2,942 images in the 2 US cities were obtained from Google Street View and images from the 2 Austria cities were collected manually on site.

The Place Pulse project found that the range of perception elicited by the images of the 2 US cities was larger than that of in the 2 Austria cities. The authors interpreted this as the evidence that the cityscapes of the 2 US cities were more

unequal than the 2 Austria cities. The authors argued that besides crime and disorder, many other dimensions, such as whether the place look lively, modern, inspiring, classy, abandoned, congested, colorful or beautiful, can be used to explore connections between aspects of urban perceptions and other social dimensions, such as entrepreneurship, civic engagement and high-school completion, among other things.

The research proved that online images can be used to create reproducible quantitative measures of urban perception and characterize the inequality of different cities.

The Streetsocore project created a computational model for perceived safety based only on image features. It used support vector regression on commonly used features extracted from the images, such as geometric texture, color histograms, and GIST. It also used the model's predicted data to create high-resolution maps of perceived safety for 21 cities in the United States, scoring around 1 million images from Google Street View. The research converted the target labels in [1] to a ranked score for each image using the Microsoft Trueskill algorithm [3]. The authors argued that an algorithm trained with a few images from a given area can be used to map a significant area of its surroundings, which is useful to know for future crowdsourced studies.

With recent developments in deep learning methodologies on computer vision, Convolutional neural network (CNN) has become the most commonly applied and state-of-the-art technology in image recognition. Therefore, our project intended to further the research of [2] with CNN. We would also specifically apply our model on a set of Toronto street view images, so as to help the City of Toronto to better understand people's perceived safety in the city.

As [2] also pointed out that future works can focus on identifying the objects and features that help explain the evaluative dimensions of a streetscape. Therefore, we also worked on model interpretation. This would further help the scholars and city planners to understand what types of elements in the living environment would be considered as safer by the residents so as for them to improve our living environment.

To achieve our research objectives, we made use of the predicted data open-sourced by [2], and created 2 groups of

deep learning classification models that could predict binary labels: More Safe and Less Safe based on Google Street View images. One group was generated from our self-designed CNN, and the other group was generated from transfer learning using pre-trained weights of ResNeXt50 developed by [4]. We used averaging and weighted ensembling of the predictions of those models to produce the final prediction.

Our models were then applied to predict on Toronto street view images. At last, we generated interpretations and explanations for our models using the Locally Interpretable Model-Agnostic Explanations (LIME) package.

LIME is a technique that explains the predictions of any machine learning classifier or regressor in an interpretable and faithful way, by learning an interpretable model locally around the prediction. It can also explain models by presenting representative individual predictions and their explanations in a non-redundant way, framing the task as a submodular optimization problem. LIME can be used to: decide if one should trust a prediction; choose between models; improve untrustworthy classifiers; and identify why a classifier should not be trusted. When explaining a prediction, LIME presents textual or visual artifacts that provide qualitative understanding of the relationship between the instance's components and the model's prediction [5].

II. METHODOLOGY

A. Data Collection

Proportion of Toronto Zoning Data. Table I was the proportion of Zoning classes in Toronto. It was provided by the City of Toronto. As we aim to build models that can predict well on the Toronto images, when we chose our Training and Test sets, we stratified sampled based on this proportion.

Boston and New York Data with Predicted Q-score.

There were 2 datasets - *streetscore_boston.csv* and *streetscore_newyorkcity.csv* open-sourced by the Streetscore project available at [6]. The two datasets both have 3 columns: latitude, longitude and q-score.

Boston Zoning data. We obtained a file *Zoning_Subdistricts.shp* from City of Boston's open data hub [7]. It contained a 'geometry' column which is geometric polygon data as well as a 'Zone_Desc' column, which was relevant to the zoning class in Table I.

Toronto data. We obtained file *ADDRESS_POINT_WGS84.shp* from the City of Toronto's Open Data Catalogue [8]. It contained 525,545 geolocation points in Toronto, but no zoning class. We randomly sampled 2,100 Toronto data to fetch Toronto images. We would use our final ensemble strategy to predict on those Toronto images.

Street View Images. We fetched all images through Google Street View Static API [9] based on the geolocation

information (latitude and longitude) provided as in the data mentioned above, and the google-streetview package [10]. Note that the API provides several parameters with which users can define while fetching images. However, it is not clear what parameter settings were adopted by the MIT Streetscore project. Therefore, we decided to set some parameters (fov, heading, pitch, radius) as default.

B. Data Exploration and Preparation

a) Creating Target Column "safety":

In the MIT Streetscore data of New York City and Boston, perceived safety scores (q-score) was continuous, ranging from -4.0 to 43.0 (with the mean of 24.86 and the median of 24.88). The distribution of the q-score is shown in Fig 1.

To create a target binary column "safety", we first combined all samples in the 2 cities and then separated the samples into 2 bins based on the distribution of the q-scores. For samples in bin 1, we defined their safety as 0, which represents Less Safe; for samples in bin 2, we defined their safety as 1, which represents More Safe. Finally, we saved the two cities' data into separate data tables: *boston_safety.csv* and *ny_safety.csv*.

b) Creating Subdistrict Variable for Boston Data:

The objective of this procedure was to make our subsampled Boston data's subdistrict classes representative to Toronto's zoning classes in the next procedure (Stratified

TABLE I. PROPORTION OF TORONTO ZONING CLASS

Zone Class	Percent (%)
residential	49.89
open space	17.86
employment industrial	13.9
unassigned	5.83
commercial	5.34
utilities	5.08
institutional	2.1




Fig. 1. Distribution of q-score of New York and Boston Streetscore data.

Downsampling). For the 2 US cities, we only found available subdistrict data of Boston that were comparable to the zone classes provided in Table I. Therefore, we decided to use only Boston data in order to meet the stratified sampling requirement.

For *streetscore boston.csv*, we used the Shapely package [11] to create a column called ‘Coordinates’, which was the geometric point generated based on the original ‘latitude’ and ‘longitude’ columns. Then we used the GeoPandas package [12] to conduct spatial join with the *Zoning_Subdistricts.shp* which has a ‘geometry’ column with geometric polygons.

Note that after the spatial join, only 23.26% (53,401 out of the 229,564) of the original Boston Data samples fell into the zones provided by *Zoning_Subdistricts.shp* and so they were left for stratified subsampling.

c) Stratified Downsampling of Boston Data:

After the ‘safety’ and ‘subdistrict’ columns were generated, we performed stratified downsampling based on both of them.

Stratified sampling based on Zone classes of Toronto.

Note that the subdistrict classes of Boston Data (Table II’s ‘subdistrict’ column) and the Zone classes of Toronto provided in Table I did not exactly match. Therefore, We considered “Miscellaneous” + “Mixed Use” of Boston subdistrict as one class, and when subsampling, we would make it stratified of the proportion of “unassigned” + “utilities” in Toronto.

As we were to select 20,000 Boston samples, we calculated for each subdistrict, how many samples we should select based on Table I, and listed it in Table III.

Stratified sampling based on target column ‘safety’.

While we sampled Boston data based on Table III, we also made sure that for each subdistrict in Boston, the proportion of ‘safety = 0’ in the target column was equal to that of ‘safety = 1’ as much as possible.

TABLE II. COUNT AND PROPORTION OF 53,401 SAMPLE’S BOSTON SUBDISTRICT CLASS

Subdistrict	Count	Percent (%)
Residential	33749	63.2
Open space	5273	9.87
Business	4650	8.71
Industrial	3452	6.46
Mixed Use	2973	5.57
Miscellaneous	1890	3.54
Comm/Instit	1414	2.65

TABLE III. THE 20,000 BOSTON SAMPLES WE AIMED TO SELECT BY SUBDISTRICT

Subdistrict	Samples should select
Residential	9978
Open space	3571
Industrial	2780
Mixed Use and Miscellaneous	2182
Business	1068
Comm/Instit	420

Note that not all subdistrict classes meet the desired quantity requirement. In subdistrict “Industrial”, there were only 505 samples which have “safety = 1”, less than the ideal 1390 “safety = 1” samples that we would like to sample. So for ‘Industrial’, more samples with safety = 0 were selected. Table IV shows the 20,000 sampled Boston data, count by both ‘subdistrict’ and ‘safety’.

d) Train / Test split:

For the 20,000 Boston samples, we split them into 80:20 ratio as Boston Train Data and Boston Test Data.

TABLE IV. SUBSAMPLED 20,000 BOSTON DATA - COUNT BY SUBDISTRICT AND SAFETY

Subdistrict	Safety	Number of Samples
residential	0	4989
residential	1	4989
open space	0	1786
open space	1	1786
business	0	534
business	1	534
Industrial	0	2275
Industrial	1	505
Mixed Use	0	663
Mixed Use	1	652
Miscellaneous	0	428
Miscellaneous	1	439
Comm/Instit	0	210
Comm/Instit	1	210

e) *Image Fetching and Pre-Processing:*

Image Fetching. Note that not all geolocation in our datasets had an image that can be fetched. For the 16,000 Boston Train Data, we were able to fetch 15,928 Boston Train Images. For the 4,000 Boston Test Data, we fetched 3,976 Boston Test Images. For the 2,100 Toronto Data, fetched 2,034 Toronto Images. Fig 2 shows some sample Boston street view images labeled as More Safe and Less Safe.

It is notable that Google periodically updates street view images, which means the street view images we fetched and then used to build and test models were mostly updated from the images used by the Streetscore project to arrive at the q-score. Table V lists the updated year of images we fetched for the Boston Train and Test Images, according to the metadata of the fetched images provided by the Google API.

Image Pre-Processing. Each downloaded street view image has a Google logo. We cropped out a small edge from the bottom and from the right so as to remove the Google logo in each image.

As we fit images into CNN, images were resized to 224*224 pixels (in transfer learning) and 100*100 (in our own CNN) pixels.

C. *Model Building*

a) *Our Self-Designed CNN Model:*

We created 2 self-designed CNN models. The first model

TABLE V. BOSTON TRAIN AND TEST IMAGES' UPDATE YEAR

Year Updated	Boston Train Image	Boston Test Image
NaN	11	2
2007	130	38
2008	6	2
2009	79	18
2010	26	6
2011	706	201
2012	23	6
2013	1853	440
2014	428	83
Sum NaN, 2007 to 2014 (Percent out of Total)	3262 (20.5%)	796 (20.0%)
2015	243	82
2016	660	174
2017	1540	387
2018	10200	2534
2019	23	3
Sum 2015 to 2019 (Percent out of Total)	12666 (79.5%)	3180 (80.0%)
Total	15928	3976




Fig. 2. Sample Boston street view images. The above 4 images were with label 'more safe' and the bottom 4 images were with label 'less safe'.

was trained and validated on the Boston Train Images and tested on Boston Test Images. This model was used to experiment with the best ensembling strategy.

The second model was trained on both the Boston Train Images and Boston Test Images. This model was used to investigate whether increasing training data could boost model performance. When we ensembled the predictions on the Toronto Images, we used the prediction of this model.

Pre-Processing: As part of pre-processing, each training image was converted to a matrix of 100*100*3. We applied one-hot encoding to the target variable to improve efficiency.

CNN layers: Self-designed CNN models were built with 2 convolutional layers, each followed by a pooling layer. The output from these layers were passed to a flatten layer and were fed to a couple of dense layers. After the initial set of runs, we realized that the model was over-fitting. Therefore, we added 2 dropout layers.

Dropout is a technique used to prevent overfitting and to approximately combine exponentially many different neural network architectures effectively. Dropout refers to temporarily removing a unit from the neural network along with its incoming and outgoing connections. As a result of Dropouts, hidden units trained with Dropout learn to work with a randomly chosen sample of other units. This makes the hidden units more robust and they do not rely on other units to correct its mistake.

After multiple runs, we found that Adam optimizer gave us the best results with default values for learning rate, decay and momentum. For the first model, we ran 20 epochs and for the second model, we ran 25 epochs.

b) *Transfer Learning:*

We chose ResNeXt50 as our pre-trained models in transfer learning. It is a simple architecture which adopts VGG/ResNets' strategy of repeating layers, while exploiting the split-transform-merge strategy in an easy and extensible way. A module in the ResNeXt network performs a set of

transformations, each on a low-dimensional embedding, whose outputs are aggregated by summation[4].

ResNeXt won the second place in image classification in the 2016 ImageNet Large Scale Visual Recognition Challenge (LSVRC) competition. ResNeXt50 has a relatively small amount of tunable parameters which allowed us to go with smaller GPU memory.

We used Adam as optimizer, and trained the models with 50 epochs and a starting learning rate of 0.0001.

Image augmentation on the training set was performed randomly on the fly during the model training process. For one particular training image, 1 augmented image was generated.

Five-fold cross-validation was implemented in the training process. For each cross-validation fold, the model with the best performance among all 50 epochs was saved. We could use the 5 submodels for ensembling. After cross-validation, the pre-trained networks were used to train models on the full training data set.

Therefore, a total of 6 models were generated (5 submodels from cross-validation + 1 model trained on full data set) in the transfer learning process.

D. Ensembling

We implemented the following ensembling methods:

Averaging Predictions on Augmented Test Images. For the 6 models produced in transfer learning, when generating predictions on the test images, for each of the test images, 5 random augmented images were generated. Those 5 predictions were given the same weight and averaged out to serve as the prediction of this particular image.

To find the best ensemble strategy, we further used 7 models (the first model produced by our self-designed CNN structure and the 6 models produced by transfer learning) to predict on the Boston Test Images, and then conducted the following ensembling methods:

Averaging Ensemble. We tried 4 averaging ensemble strategies on the 7 models: averaging prediction of all single models; averaging predictions of models with best accuracy; averaging predictions of models best at predicting safety = 0 and safety = 1 respectively.

Averaging + Conditional Ensemble. We tried 2 conditional ensemble strategies: the first one is when predicting safety = 0, we used the predictions best at predicting safety = 0 produced by averaging ensemble, else we used the prediction by predictions best at predicting safety = 1 produced by averaging ensemble. The second one is when predicting safety = 1, we use the predictions best at predicting safety = 1 produced by averaging ensemble, else we use the predictions best at predicting safety = 0 produced by averaging ensemble.

Averaging + Weighted Ensemble. We tried 3 weighted ensemble strategies with different weight pairs to averaged predictions best at predicting safety = 0 and safety = 1 respectively.

Final Ensemble Prediction Strategy. We compared the performance of ensembled predictions and picked the one ensemble strategy that had a good overall accuracy and was good at predicting of both targets.

E. Model Evaluation

For each of the single model predictions and the ensembled predictions on the Boston Test Images, we calculated the confusion matrix so as to evaluate model performance as well as to navigate us towards the next and the finally the best ensembling strategy.

F. Predictions on Toronto Images

We applied the final ensemble prediction strategy on the predictions of the 2,034 Toronto images.

G. Model Interpretation

At last, we applied LIME on a number of randomly chosen Toronto street view images predicted by our models and conducted qualitative analysis on the results.

LIME can generate a local interpretation for a particular image, which will in turn help us to understand all the factors that contribute towards or against a prediction. It attempts to understand the model by perturbing the input of data samples and understanding how the predictions change. It creates variations of the image by segmenting image into superpixels and turning superpixels off or on. Superpixels are interconnected pixels with similar colors.

III. RESULTS

A. Performance of Single Models on Boston Images

The loss and accuracy of single models produced by our own CNN model and transfer learning are shown in Table VI. Most of our models experienced overfitting, except for the one that produced by our own CNN model on both Boston Train and Test Images. Our single models' validation accuracy was all around 70%.

For the 2 models produced by our own CNN structure, the model trained on both Boston Train and Test Images had better performance and less overfitting comparing the model trained by the same CNN structure on Boston Train Images only.

B. Confusion Matrix and Performance of Single and Ensembled Models on Boston Test Image.

The confusion matrix and performance of single models and all ensembled models on Boston Test Images are shown in Table VII.

Single Models. Among the single models, 2 models - cv3 and full were better at predicting safety = 0; the other models -

cv0, cv1, cv2 cv4, own were better at predicting safety = 1. For the 6 models produced by transfer learning, their accuracy on the Boston Test Images was better than the validation accuracy on Boston Train Images.

Averaging Ensemble. Averaged predictions had better accuracy than single models.

Overall, ensembling of ‘all models’ yielded the best accuracy (73.9%). However, this ensembled model is neither best when predicting safety = 0 along nor best at predicting safety = 1 alone.

Model ‘cv3full’ was best at predicting safety = 0 (TN_Rate 0.756619) and model ‘cv0cv1cv2cv4own’ was best at predicting safety = 1 (TP_Rate 0.779363), but they were weak at predicting the other target respectively. We would use these 2 models in conditional and weighted ensembling.

Averaging + Conditional Ensemble. We tried 2 conditional ensembling:

Conditional ensemble1 - When predicting safety = 0, we use the prediction by ‘cv3full’; else we use the prediction by ‘cv0cv1cv2cv4own’.

Conditional ensemble2: - When predicting safety = 1, we use the prediction by ‘cv0cv1cv2cv4own’; else we use the prediction by ‘cv3full’.

Confusion Matrix showed that show that ‘conditional ensemble1’ model was good at prediction target 0, at the expense of prediction 1. On the contrary, ‘conditional ensemble2’ model was good at prediction target 1, at the expense of prediction 0. Therefore we further tried weighted ensemble.

Averaging + Weighted Ensemble. We tried 3 weighted ensembling with different weight pairs to model ‘cv3full’ and model ‘cv0cv1cv2cv4own’, as detailed in Table VII.

Comparing results throughout Table VII, we found that the ‘weighted ensemble2’ had an accuracy of 74.1%, and was good at predicting both safety = 0 and safety = 1.

Therefore, we would consider the method used in ‘weighted ensemble2’ as our final ensembling strategy.

C. Predictions on Toronto Street Views

We applied the ensemble strategy the same way as ‘weighted ensemble2’ on the 2,034 Toronto images. As a result, 1,031 of the images were predicted as 0 (less safe), and 1,003 were predicted as 1 (More Safe). Fig 4 is a map visualization of the predictions. It shows that in the downtown area, there are more street views predicted as More Safe comparing to the other areas.

D. LIME

Our qualitative analyses on LIME’s interpretation of our models’ predictions on some randomly chosen Toronto street view images yielded useful insights. We found that plants,

trees, lawn, newer looking and brightly colored houses seemed to heavily contributed towards a More Safe perception. On the contrary, electrical wires hanging over houses, dull colored and older houses, shadows, and crack on the ground contributed towards Less Safe perception.

Fig. 3 shows 4 images and their LIME interpretations. Among them, image (a) and image (b) were predicted as More Safe by our models; image (c) and image (d) were predicted as Less Safe by our models. LIME highlighted the parts of images that contributed towards a prediction in green color (pros) and the parts of the image that contributed against a prediction in red color (cons).

In all of the 4 images, the trees contributed towards a More Safe perception, whereas dull colored and older houses, and shadows contributed towards a Less Safe perception. In image (a), LIME seemed to interpret that crack on the ground at the right corner of the image contributed to the Less Safe prediction. In image (c), LIME interpreted that wires running overhead the house made the model to predict the image as Less Safe.

IV. DISCUSSIONS

A. Transfer Learning vs Self-Designed CNN

Models generated by transfer learning and our self-designed CNN models performed similarly. But Transfer learning using ImageNet pretrained weights costed much higher computational resources. Previous research [13] challenged the conventional wisdom of ImageNet pre-training for dependent task. Their research found that in in object detection and instance segmentation tasks on the *COCO* dataset, standard models trained from random initialization performed no worse on ImageNet pre-training. Our research further demonstrated that in image classification, standard models also performed no worse on ImageNet pre-training.



Fig. 3. Examples of model interpretation by LIME. Image (a) and image (b) were predicted as More Safe by our models; Image (c) and image (d) were predicted as Less Safe by our models.

TABLE VI. PERFORMANCE OF SINGLE MODELS TRAINED FROM TRANSFER LEARNING AND OUR OWN CNN ON BOSTON IMAGES

	Transfer Learning from ResNeXt50								Own CNN	
Training data	CV0 of Training	CV1 of Training	CV2 of Training	CV3 of Training	CV4 of Training	mean of 5 CVs	std of 5 CVs	Full Training	Full Training	Full Training + Test
train_loss	0.167847	0.070979	0.085303	0.051624	0.052649	0.085680	0.042939	0.058055	0.3158	0.4645
train_acc	0.947100	0.972844	0.970884	0.980929	0.979752	0.970302	0.012226	0.978464	0.8696	0.7783
val-loss	1.802394	1.540489	1.852702	1.808886	1.562203	1.713335	0.133567	-	0.6401	0.5703
val-acc	0.669805	0.672003	0.691994	0.672214	0.692622	0.679727	0.010308	-	0.6800	0.7270
Name in Ensemble	CV0	CV1	CV2	CV3	CV4	-	-	Full	Own	-

TABLE VII. CONFUSION MATRIX AND PERFORMANCE OF MODELS ON BOSTON TEST IMAGES

Ensemble Type	Ensemble Strategy	Model Name	FN	FP	TN	TP	TN_Rate	TP_Rate (Recall)	Accuracy	F1 score	Precision
Single Model	NA	cv0	485	704	1449	1337	0.673014	0.733809	0.700881	0.692208	0.655071
		cv1	442	775	1378	1380	0.640037	0.757409	0.693836	0.693990	0.640371
		cv2	537	649	1504	1285	0.698560	0.705269	0.701635	0.684239	0.664426
		cv3	653	558	1595	1169	0.740827	0.641603	0.695346	0.658777	0.676896
		cv4	505	684	1469	1317	0.682304	0.722832	0.700881	0.688988	0.658171
		full	579	536	1617	1243	0.751045	0.682217	0.719497	0.690364	0.698707
		own	552	742	1412	1270	0.655829	0.697036	0.674717	0.662666	0.631527
Averaging Ensemble	All Models	all models	435	602	1551	1387	0.720390	0.761251	0.739119	0.727893	0.697335
	Models with Best Accuracy	cv0cv2cv4full	462	620	1533	1360	0.712030	0.746432	0.727799	0.715413	0.686869
	Models Best at Predicting Safety = 0	cv3full	576	524	1629	1246	0.756619	0.683864	0.723270	0.693764	0.703955
	Models Best at Predicting Safety = 1	cv0cv1cv2cv4own	402	664	1489	1420	0.691593	0.779363	0.731824	0.727087	0.681382
Averaging + Conditional Ensemble	For safety = 0, use cv3full; else use cv0cv1cv2cv4own	conditional ensemble1	650	430	1723	1172	0.800279	0.643249	0.728302	0.684579	0.731586
	For safety = 1, use cv0cv1cv2cv4own; else use cv3full	conditional ensemble2	328	758	1395	1494	0.647933	0.819978	0.726792	0.733432	0.663410
Averaging + Weighted Ensemble	cv3full * 0.4; cv0cv1cv2cv4own * 0.6	weighted ensemble1	455	574	1579	1367	0.733395	0.750274	0.741132	0.726548	0.704276
	cv3full * 0.45; cv0cv1cv2cv4own * 0.55	weighted ensemble2	460	570	1583	1362	0.735253	0.747530	0.740881	0.725626	0.704969
	cv3full * 0.55; cv0cv1cv2cv4own * 0.45	weighted ensemble3	490	563	1590	1332	0.738504	0.731065	0.735094	0.716707	0.702902

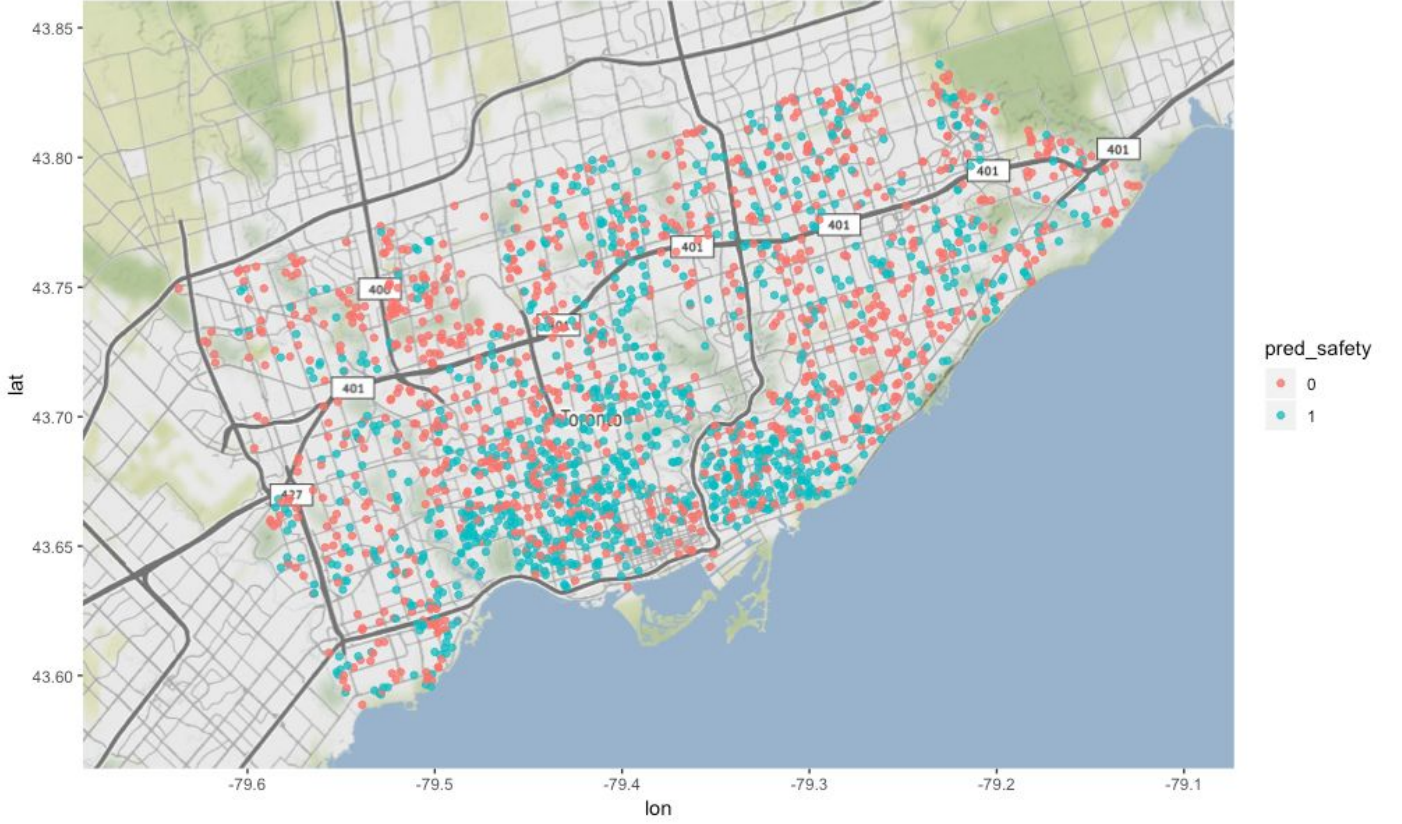


Fig. 4. Map of our models' final predictions of perceived safety on street view of 2,034 Toronto geolocations

B. Ensemble Learning

Averaging Predictions on Augmented Test Images. For the 6 models produced by transfer learning, their accuracy on the Boston Test Images (Table VII) was better than the validation accuracy on the Boston Train Images (Table VI). This proved that generating more augmented images on the test set and then average out the prediction could boost model accuracy.

Ensembling. Our research showed the effectiveness of ensembling: the accuracy of all ensembling models was better than that of single models. Our final ensemble strategy boosted the prediction accuracy to 74.1% from ~70% accuracy by single models. Confusion matrix of each ensembled models proved to be an effective tool in guiding us to find the best ensembling strategy.

C. Model Performance

Despite the effectiveness of ensembling, our single models' performance was not satisfactory. We assume there may be several reasons:

First, ~80% of the images have been updated from the time when the Streetscore project (published in 2014) generated the q-score based on the street view images available through

Google Street View Static API at that time..As city scenery changes rapidly, we argue that the images for the same locations might have changed a lot during the past 4 years. Also it is not clear what parameters the Streetscore project used when fetching images using the Google API. For a same location, the different parameter settings would get very different images.

Second, the number of training images we used was not large enough. We used only 15,928 Boston Train Images, for the submodels saved during cross-validation, the training images are even fewer (~12,700). The fact that when applying our own CNN structure on both the Boston Train Images and the Boston Test Images, the model accuracy (0.7270) was much higher than that of the model trained on Boston Train Images alone (0.6800) proved our this argument.

Fig. 5 shows some of the extremely wrongly predicted Boston Test Images based on our final predictions. The upper 4 images were predicted as having an over 0.993 probability of being Less Safe, but the real label should be More Safe. The below 4 images were predicted as having an over 0.991 probability of being More Safety, but the real label should be




Fig. 5. Failure cases. The upper images are top 4 wrongly predicted as safety = 0; The below images are top 4 wrongly predicted as safety = 1.

Less Safe. It seemed that our models tended to predict images with empty spaces and gloomy color to have a high probability of being Less Safe, and images with trees to have a high probability of being More Safe.

The failure cases echoed what we had observed in qualitative analyses of the LIME results, and double confirmed that our models gave too many weights to trees in the predictions of More Safe; and gave too many weights to dull colored and older houses in the prediction of Less Safe. These may have contributed to the model overfitting. We believe that this issue could be improved if more training images will be used in the future.

Third, information shifting when generating the target label. The Streetscore project used the PlacePulse1.0 dataset provided by [1] for prediction. The original PlacePulse1.0 dataset had 3 continuous target labels: QS Safer, QS Unique and QS Upperclass ranging from 1-10. In [1], the author noted that the mapping between images and locations is not one-to-one. Therefore, for a large number of locations, they captured more than one image, by pointing the camera in two or more directions.

The Streetscore project transformed the PlacePulse1.0 dataset to create one continuous label - q-score, ranging from 1-10. It is notable that our project used data *predicted* by the Streetscore project. In Streetscore project's open-sourced datasets, the label was a continuous q-score ranged from -4.0 to 43.0. And we further generated a binary label 'safety' based on the predicted q-score. This series of information shifting may have contributed to the model inaccuracy as well.

Fourthly, we could try more aggressive image augmentation setups.

D. Model Interpretation

Our research showed that model interpretation tools are effective in understanding the models' predictions. Qualitative analysis of model interpretation results could provide profound insights for researchers and city planners to understand what elements are mostly perceived as More Safe by residents. This makes it possible for the city planners to

improve our living environment, as well as to generate *predicted* quantitative data for psychologists to study people's perceptions towards the living environment.

V. LIMITATIONS AND FUTURE WORKS

Limitations. Our research had several limitations, mainly about the training data.

First, the purpose of this project is to predict the perceived safety for Toronto, however, we do not have labeled training images for Toronto. The performance of our model would not be as robust as it would be when using Toronto based training data. It is also not possible to effectively evaluate our predictions on Toronto images.

Second, the target label we generated was based on predicted values.

Third, there was a time gap between the time we fetched the images and the time the images were fetched by the Streetscore project (2014) to predict the q-score. This would cause model inaccuracy.

Future Works. Firstly, we suggest future research to collect manually-coded safety score towards Toronto street view images, and shorten the time gap between image fetching for model generation and image fetching for model prediction.

Secondly, future works could analyze the correlation between predictions on Toronto street safety and social factors such as crime and income, like [1] and [2] did for other cities.

Thirdly, it would be interesting to try whether different image augmentation setups and different pre-trained models could boost model performance.

Fourthly, we have just conducted qualitative analyses on some randomly selected LIME results. It would be very exciting to see systematic qualitative research using large scale LIME predictions.

VI. CONCLUSIONS

In this project, we conducted exploratory research by using a predicted dataset open-sourced by the Streetscore project and by applying CNN, transfer learning, ensemble learning, and model interpretation techniques to predict people's perceived safety on Toronto street views.

Our research laid a foundation for the City of Toronto to develop an automated predictor that can predict city-scene at scale for the entire city, so as to provide the residents and urban planners with this information for all parts of the city.

Our research also suggested that qualitative analyses on a large number of model interpretation results on the models' prediction on perceived street safety could help researchers and city planners to understand what elements are mostly perceived as More Safe by residents and thus improve people's living environment.

ACKNOWLEDGMENT

We would like to thank the kind guidance by Faezeh Khabbaz, our course instructor of ML1030: Machine Learning at Scale; Matthew Tenney, Data Science Supervisor of Data Analytics and Visualization team in Geospatial Competency Centre of the City of Toronto; and Ryan Garnett, Manager Geospatial Data Integration & Access of City of Toronto.

REFERENCES

- [1] Salesses, P., Schechtner, K., & Hidalgo, C. A. (2013). The collaborative image of the city: mapping the inequality of urban perception. *PloS one*, 8(7), e68400.
- [2] Naik, N., Philipoom, J., Raskar, R., & Hidalgo, C. (2014). Streetscore-predicting the perceived safety of one million streetscapes. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops* (pp. 779-785).
- [3] Herbrich, R., Minka, T., & Graepel, T. (2007). TrueSkill™: a Bayesian skill rating system. In *Advances in neural information processing systems* (pp. 569-576).
- [4] Xie, S., Girshick, R., Dollár, P., Tu, Z., & He, K. (2017). Aggregated residual transformations for deep neural networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 1492-1500).
- [5] Ribeiro, M. T., Singh, S., & Guestrin, C. (2016, August). Why should i trust you?: Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining* (pp. 1135-1144). ACM.
- [6] <http://streetscore.media.mit.edu/data.html>
- [7] https://bostonopendata-boston.opendata.arcgis.com/datasets/b601516d0af44d1c9c7695571a7dca80_0
- [8] <https://www.toronto.ca/city-government/data-research-maps/open-data/open-data-catalogue/#f71a13c4-fb51-6116-57b7-1f51a8190585>
- [9] <https://developers.google.com/maps/documentation/streetview/intro>
- [10] <https://pypi.org/project/google-streetview/>
- [11] <https://pypi.org/project/Shapely/>
- [12] <http://geopandas.org/>
- [13] He, K., Girshick, R., & Dollár, P. (2018). Rethinking imagenet pre-training. *arXiv preprint arXiv:1811.08883*.